



Preemptive Hardware Multitasking in ReconOS

Markus Happe, Andreas Traber, Ariane Trammell

Communication Systems Group, ETH Zürich, Switzerland

ARC'15 Symposium

Bochum, 15–17.04.2015

ETH zürich

Adding preemptive hardware multitasking to a reconfigurable system-on-chip (rSoC).




Preemptive hardware multitasking
multiple tasks execute on a single slot

enable rSoCs to use smart schedulers
fairness, responsiveness, avoid starvation

Unmodified hardware source code
task preemption is transparent to developer

For context restoration, we can use task- or FPGA-specific methods.



How to switch between tasks A & B?

- (1) stop task A, save context A
- (2) reconfigure task B, restore context B, run B


What is the context?

content of all state-holding elements
i.e. flip-flops, BRAMs, LUT-RAMs

Which approaches exist?

- (i) task-specific preemption (scan-chains)
- (ii) FPGA-specific bitstream readback

Task-specific preemption techniques extract/insert only the actual context.



Add scan-chains to hardware tasks
extract/insert context from/to scan-chain


Drawbacks

modifications at source code- or netlist-level
considerable hardware overhead

Advantages

approach independent of FPGA architecture
only captures/restores used resources

Bitstream readback methods read/ write configuration of the entire slot.



Read back bitstream of a partial region
using internal configuration access port (ICAP)
read contents of all flip-flops, BRAMs, etc.


Drawbacks

bitstream format tailored to FPGA family
time-inefficient (capture/restore all resources)

Advantages

no modifications at task-level
no hardware overhead

Preemptive Hardware Multitasking in ReconOS



Methodology

bitstream readback

ReconOS Integration

ICAP controller, scheduler

Virtex-6 Evaluation

reconfiguration overhead

Our methodology is based on Xilinx PlanAhead PR toolflow.

full bitstream

system_A.bit

Generate bitstreams with PlanAhead
reset_after_reconfiguration attribute: 1

partial bitstreams

partial_A.bit

Partial bitstreams: in external memory
e.g. on compact flash card

partial_B.bit

Our methodology is based on Xilinx PlanAhead PR toolflow.

full bitstream

system_A.bit

Generate bitstreams with PlanAhead
reset_after_reconfiguration attribute: 1

partial bitstreams

partial_A.bit

partial_A_captured.bit

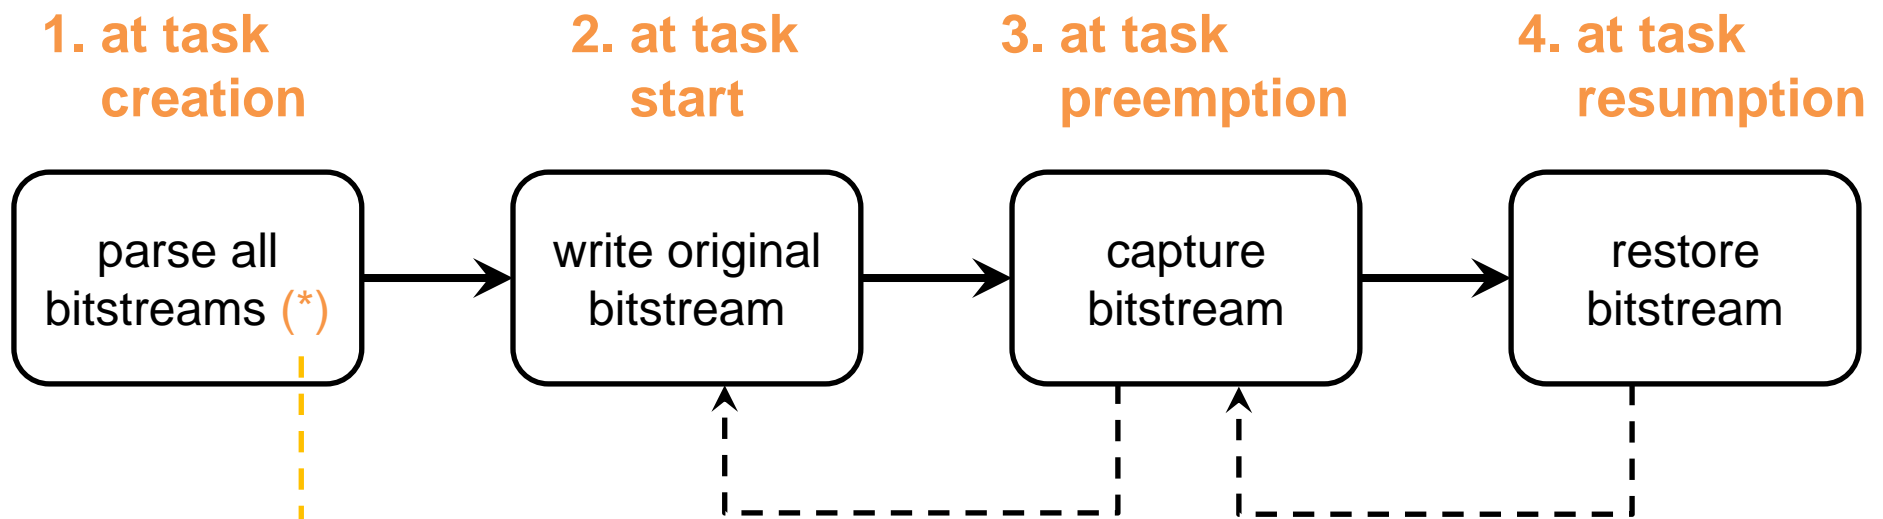
partial_B.bit

partial_B_captured.bit

Partial bitstreams: in external memory
e.g. on compact flash card

Working copy of partial bitstreams
also contain the context of preempted task
initially, working copy identical to original file


Our approach initially parses bitstreams to extract metadata about configuration frames.



(*) one bitstream per hardware slot is sufficient

Xilinx bitstreams store three types of configuration frames.

partial bitstream A



Type 1: FFs, LUTs, etc.

CLBs, input/output blocks, clocks, ...

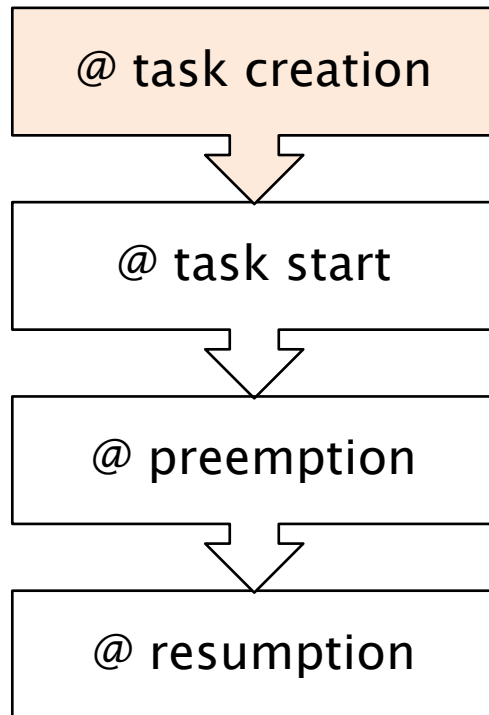
Type 2: BRAMs

configuration of BRAM contents

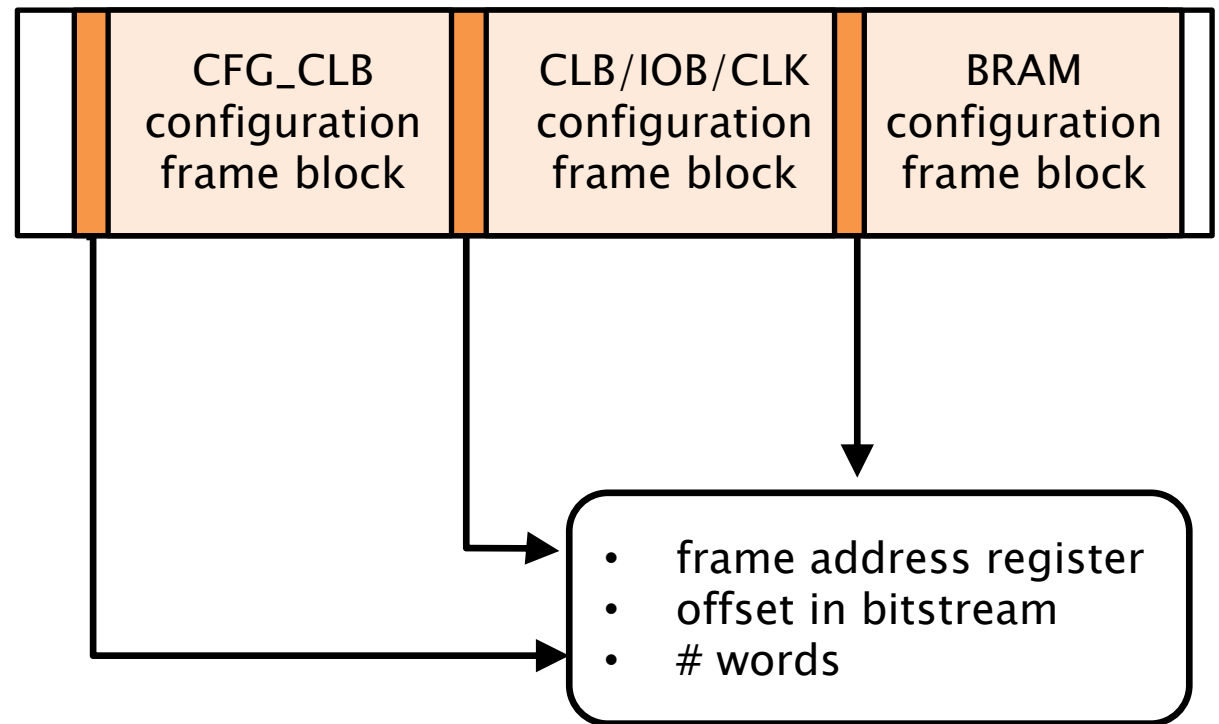
Type 3: CFG_CLB

selects which part of the FPGA is reconfigured
only present if «reset_after_reconfiguration=1»

We extract a list of metadata for all configuration frame blocks.




partial bitstream



metadata

At task start, we configure the original bitstream using ICAP write commands.



No context at task start

configure original bitstream from PlanAhead

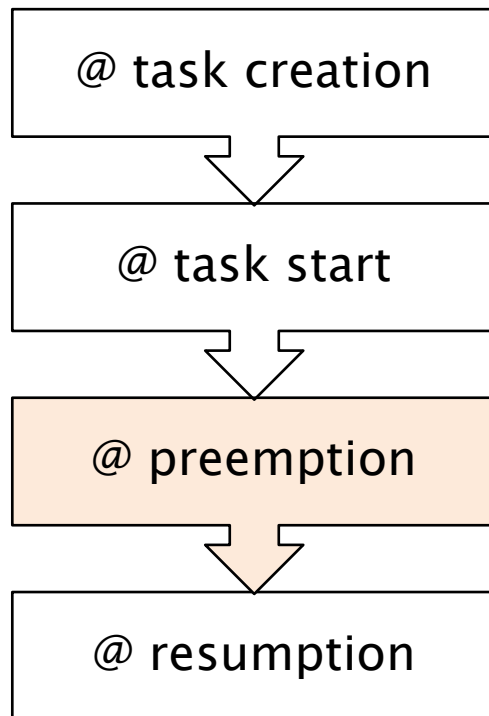
Configuration

using ICAP write commands

During reconfiguration

reset interface to/from hardware task

At task preemption, we need to capture the flip-flop states of the slot to hidden registers.



1. Prepare hardware slot

deactivate clock

2. Constrain capture command to slot

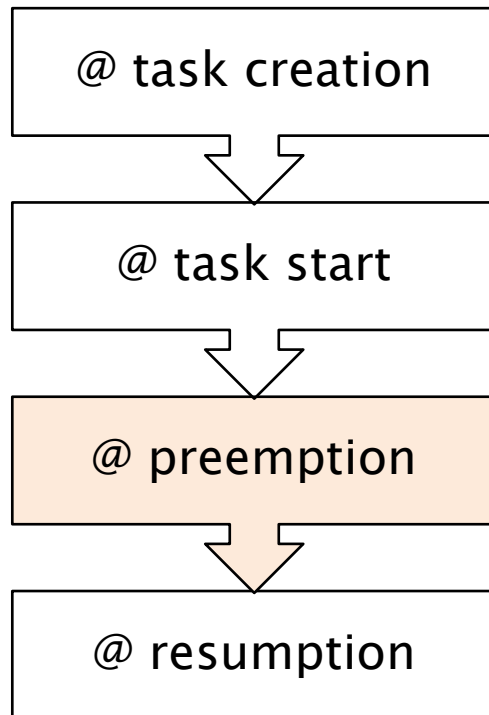
write CFG_CLB block over ICAP

3. Capture states of flip-flops

send gcapture command over ICAP

states captured to hidden registers INT0/INT1

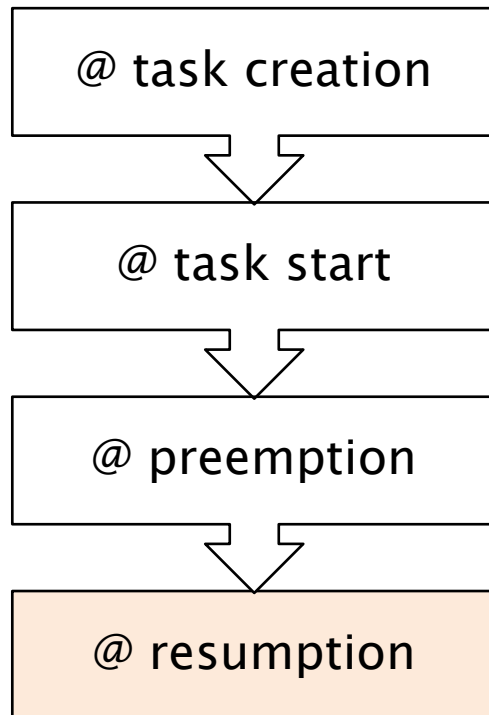
After reading back all frames, we modify certain configuration bits for the BRAMs.



4. Read-back configuration frames
as defined by metadata (except CFG_CLB block)
we update the working copy for these blocks

5. Modify BRAM configuration bits
otherwise BRAM contents will not be restored
not described in official Xilinx documentation

At task resumption, we write the modified captured bitstream to the hardware slot.



1. Prepare hardware slot

disable clock, reset slot and interfaces

2. Download captured bitstream


using ICAP write commands

3. Trigger global/set reset (GSR)

manually trigger GSR port of startup_virtex6
grestore command in bitstream does not work

4. Resume clock and hardware thread

Preemptive Hardware Multitasking in ReconOS



Methodology

bitstream readback

ReconOS Integration

ICAP controller, scheduler

Virtex-6 Evaluation

reconfiguration overhead

ReconOS extends multithreaded programming model to reconfigurable hardware.



www.reconos.de

Autonomous hardware threads

similar to POSIX software threads

access shared memory, semaphores, mqueue

Software delegate threads

represent hardware threads

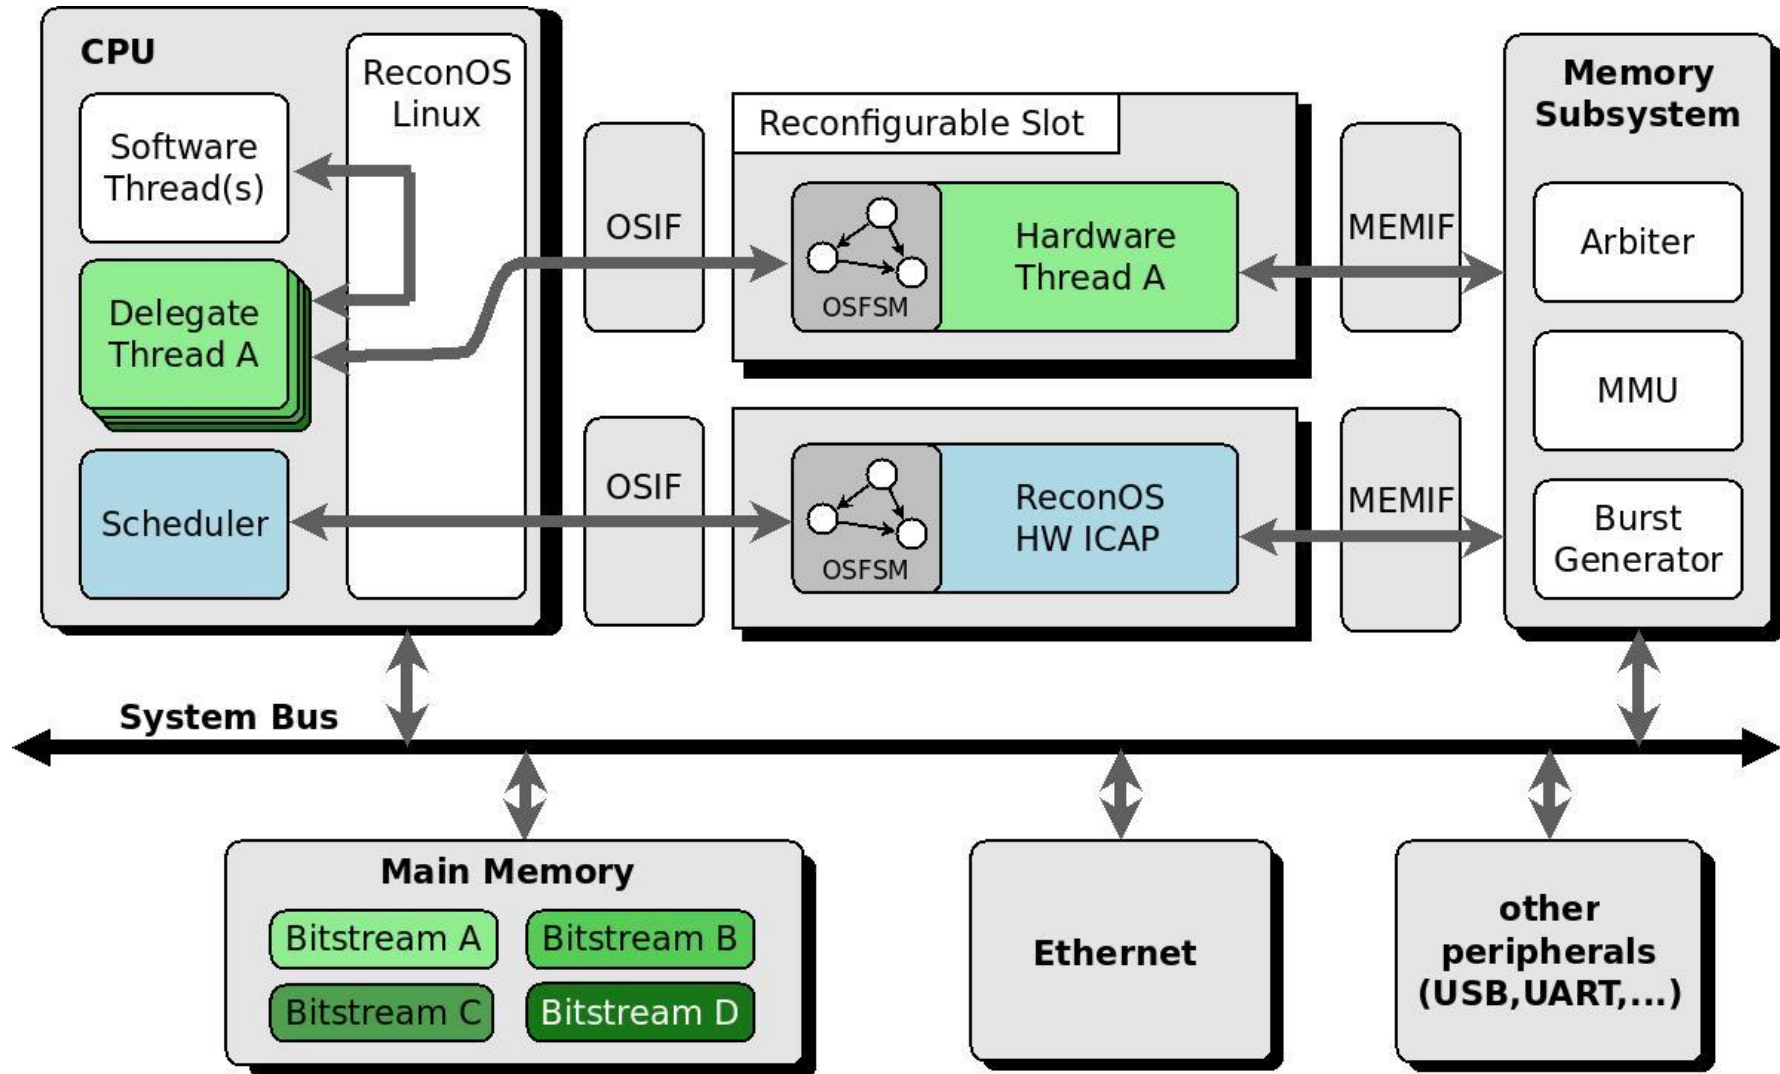
access OS services on behalf of hw threads

ReconOS @ Linux


extensions implemented in user-space

support for virtual memory

Our ReconOS Architecture has one slot and one ICAP core.



Preemptive Hardware Multitasking in ReconOS



Methodology

bitstream readback

ReconOS Integration

ICAP controller, scheduler

Virtex-6 Evaluation

reconfiguration overhead

We have performed our experiments on a Xilinx Virtex-6 FPGA for two slot sizes.



Xilinx Virtex-6 ML605 board
implements ReconOS architecture
300k flip-flops, 0.4k BRAMs (36 Kbit)

Tested slot sizes
2% of FPGA area: 361 KB bitstreams
6% of FPGA area: 741 KB bitstreams

Four hardware threads
ADD, SUB, MUL, LFSR

We implemented four hardware threads to validate the context restoration.

$$R_3 = R_1 \pm R_2$$

ADD/SUB-Threads

add/substract sw-defined registers

$$R_3 = \sum_1^{R_1} R_2$$

MUL-Thread


multiplies sw-defined registers

validate cycle-true restoration of flip-flops

LFSR-Thread

stores several LFSRs in BRAMs, shifts one LFSR

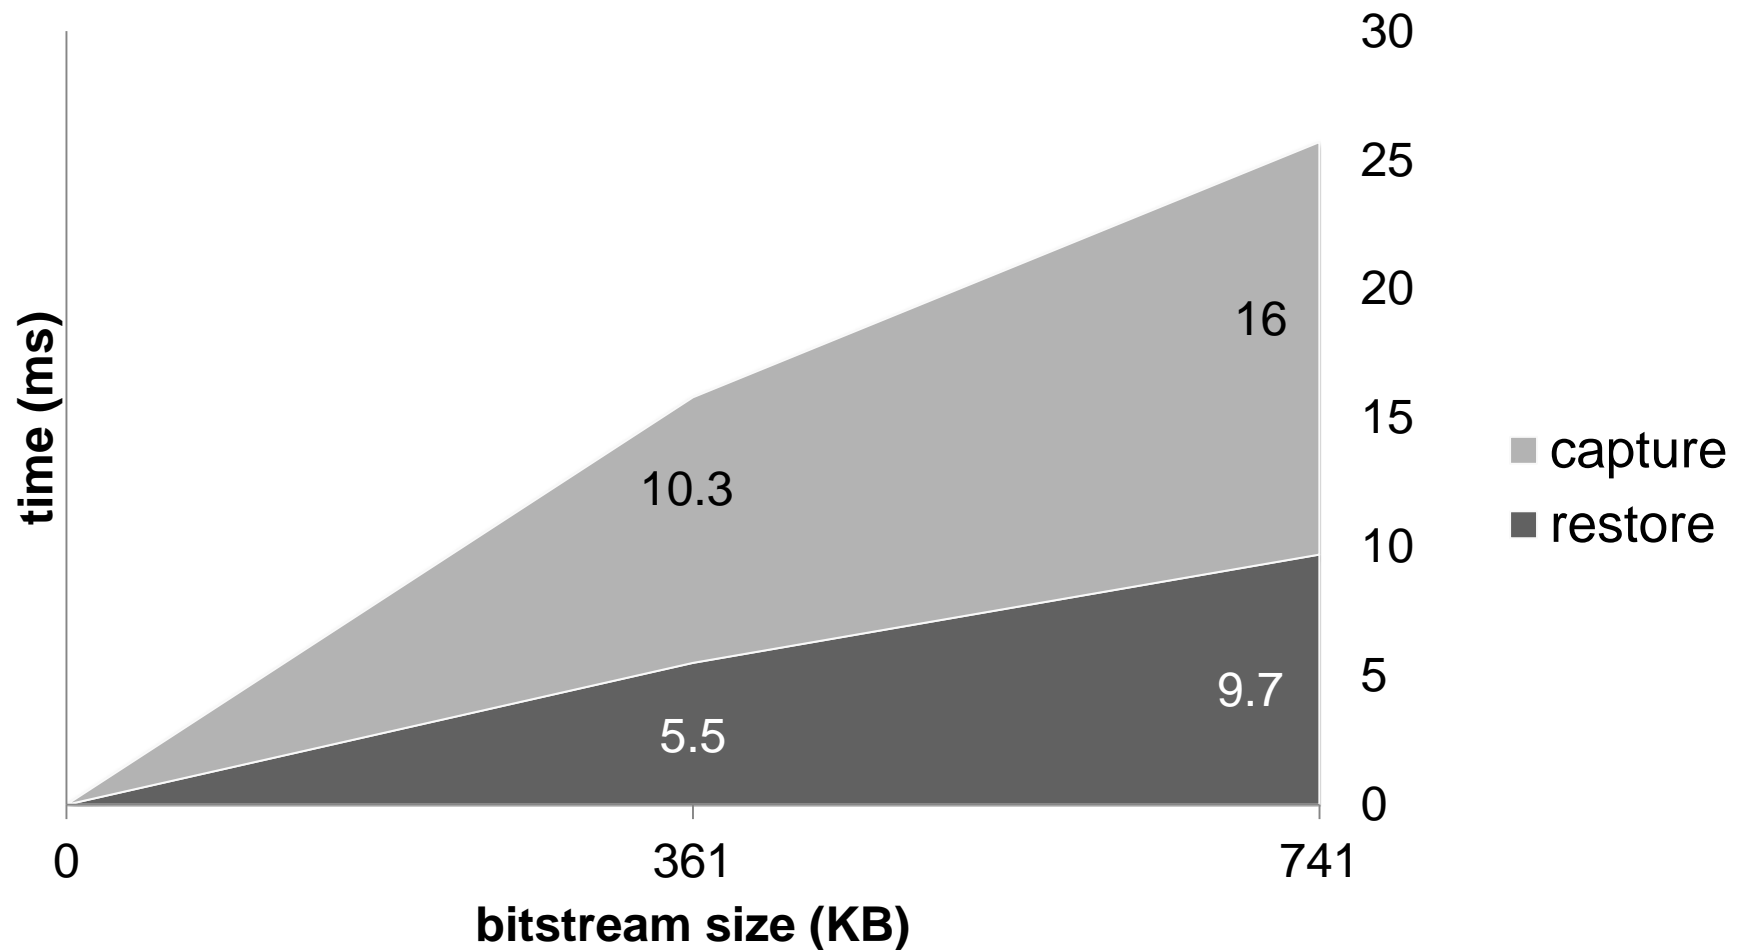
validate cycle-true restoration of BRAMs



The hardware threads only occupy a fraction of the slots, but all FFs/BRAMs are restored.

| component | #FFs | #LUTs | #BRAMs |
|---------------|------|-------|--------|
| HW slot (~2%) | 5.8k | 2.9k | 8 |
| HW slot (~6%) | 17k | 8.6k | 32 |
| ADD/SUB | 0.4k | 0.6k | 2 |
| MUL | 0.5k | 0.8k | 2 |
| LFSR | 0.5k | 0.8k | 2 |
| ReconOS ICAP | 0.3k | 0.7k | 2 |
| XPS_HWICAP | 0.8k | 0.8k | 1 |

We support context switches
in the order of milliseconds.



The maximum number of context switches per second depends on the slot size.

Max. context switches / second

361KB: 63 switches/s

741KB: 38 switches/s

Capture/restore bandwidth

361KB: 22 MB/s


741KB: 28 MB/s

Performance bottleneck

ReconOS memory interface

Periodic tasks can be executed using preemptive scheduling algorithms.

Example measurement with fixed periodic schedule



In summary, we integrated preemptive hardware multitasking to ReconOS.

Novel methodology

preemptive hardware multitasking on Virtex-6
restore states of flip-flops and BRAMs

ReconOS integration

customized ICAP controller, sw scheduler

Experimental results

achieved bandwidth: 22–28 MB/s
multiple task swaps per second

In future work, we want to support more FPGA families and relocate tasks.

Restore further FPGA resources

LUT-RAMs, DSP blocks


Support more FPGA families

e.g. Virtex-7, Zynq

Task relocation

migrate threads between slots

Preemptive Hardware Multitasking in ReconOS



Markus Happe

mhappe@ethz.ch

ETH Zürich

ARC'15, Bochum

15th April 2015